EVALUATION REPORT: ROUTER AUTO-SELECTION VS HEAVY-ONLY LLM

Generated: 2026-10-23T15:07:53.790Z | Run ID: 2026-10-23T15-07-49-348Z | Tests: 5

---

1. EXECUTIVE SUMMARY & EFFICIENCY VERDICT

Verdict: Outstanding Efficiency: Router saved 78.8% inference latency and offloaded 100% of queries away from heavy VRAM.

Efficiency Score: 87.3 / 100

Heavy Baseline Model: llama3.1:8b

2. LATENCY & THROUGHPUT COMPARISON

â¢ Router Mean Latency: 3.35 s (p50: 3.02 s)

â¢ Heavy-Only Mean Latency: 15.8 s (p50: 7.24 s)

â¢ Latency Reduction: 78.8% efficiency gain

â¢ Average Tokens/Sec: Router 130.2 tok/s vs Heavy 14.4 tok/s

3. TOKEN USAGE & DIFFERENCE

â¢ Router Tokens: 595 total (177 prompt + 418 completion)

â¢ Heavy Tokens: 656 total (82 prompt + 574 completion)

â¢ Token Usage Difference: Router used 61 fewer tokens (9.3% difference)

4. ROUTER TIER DISTRIBUTION

â¢ Light Tier: 5 (100%)

â¢ Code Tier: 0 (0%)

â¢ Heavy Tier: 0 (0%)

5. DETAILED TEST CASE BREAKDOWN (Top Cases)

Test ID	Router Tier	R-Lat(ms)	H-Lat(ms)	Saved(ms)	Eff %	Tok Diff
eval-001	light	3627	20776	17150	83%	7
eval-002	light	5086	39684	34599	87%	95
eval-003	light	3015	7236	4221	58%	-14
eval-004	light	2480	5199	2719	52%	-19
eval-005	light	2556	6161	3605	59%	-8