LLM Router - Benchmark Evaluation Report

Generated: 2026-10-18T14:16:40.115Z | Run ID: 2026-10-18T12-21-43-083Z

---

1. EXECUTIVE SUMMARY & VERDICT

Verdict: Router reduced mean latency by 21.5% vs heavy-only (KPI-2 target of 30% lower not met).

Hardware: linux x64; AMD Ryzen 7 7435HS; 24.9 GB RAM; GPU: not detected

Total Prompts: 72 labeled test prompts

Date Executed: 10/18/2026, 6:57:00 PM

2. KEY PERFORMANCE INDICATORS (KPIS)

â¢ KPI-1 Routing Accuracy: 100% (Target: 85% | Pass: YES)

â¢ KPI-2 Latency Reduction: 21.5% (Target: 30% | Pass: NO)

â¢ KPI-3 Cache Hit Rate (Run C): 58%

3. DETAILED RUN METRICS SUMMARY

Metric	Run A (Heavy Only)	Run B (Router)	Run C (Router+Cache
Mean Latency	21.6 s	16.9 s	7.08 s
p50 Latency	22.1 s	9.94 s	45 ms
p95 Latency	44.7 s	47.5 s	37.3 s
TTFT Mean	161 ms	433 ms	114 ms
Tokens / Sec	20.6 tok/s	63.7 tok/s	63.5 tok/s
Model Swaps	72	72	60
Cache Hit Rate	0%	0%	58%
Routing Accuracy	-	100%	100%
Total Wall Time	1552.6 s	1218.9 s	1080.1 s

4. ROUTING CONFUSION MATRIX (Run B)

Expected \ Routed	light	code	heavy
light	24	0	0
code	0	24	0
heavy	0	0	24

5. METHODOLOGY & PROVENANCE

Backend: Ollama local inference. Order: deterministic prompt-file order; A = forced heavy/no-cache, B = auto/no-cache, C = auto/cache-on over two full passes. Warm-up: before each run, one unmeasured forced-heavy request is completed; identical prompt and sampling settings are used. Cache is flushed before each run; for C it remains enabled for both passes. Cache results include labeled duplicate prompts. Latency is measured at the gateway from the successful request attempt through terminal SSE event; rate-limit backoff is excluded from per-request latency but included in total wall time. Router TTFT/tokens-per-second are reported from final usage metadata. Swap count is a proxy for non-cache-hit inference requests with positive final-chunk load_ms; it

is not an independent hardware swap counter. Routing accuracy and confusion matrix use the dry-run preview endpoint and expected tiers in prompts.json; this is not answer-quality evaluation. GPU/VRAM values are omitted when the stats contract reports them as unavailable. Mock-backend measurements are for smoke testing only and are not hardware performance claims.